

JavaScript is disabled on your browser. Please enable JavaScript to use all the features on this page. [Skip to main content](#)[Skip to article](#)

ScienceDirect

* [Journals & Books](#)

* [Help](#)

* [Search](#)

Gergo Gyori

IT University of Copenhagen

* [View **PDF**](#)

* [Download full issue](#)

[Search ScienceDirect](#)

[Outline](#)

1. [Highlights](#)

2. [Abstract](#)

3. [4. Keywords](#)

5. [1\ Introduction](#)

6. [2\ Related work](#)

7. [3\ Hard negative generation](#)

8. [4\ Radial metric learning](#)

9. [5\ Numerical experiments](#)

10. [6\ Conclusions](#)

11. [References](#)

12. [Vitae](#)

[Show full outline](#)

[Cited by \(86\)](#)

[Figures \(13\)](#)

1. 2. 3. 4. 5. 6.

[Show 7 more figures](#)

[Tables \(8\)](#)

1. [Algorithm 1](#)

2. [Algorithm 2](#)

3. [Table 1](#)

4. [Table 2](#)

5. [Table 3](#)

6. [Table 4](#)

[Show all tables](#)

[Pattern Recognition](#)

Volume 88, April 2019, Pages 1-12

[Hard negative generation for identity-disentangled facial expression recognition](#)

[Author links open overlay panel](#)Xiaofeng Liu [a b c](#), B.V.K. Vijaya Kumar [c d](#), Ping Jia [a b](#), Jane You [a b e](#)

[Show more](#)

[Outline](#)

[Add to Mendeley](#)

[Share](#)

[Cite](#)

<https://doi.org/10.1016/j.patcog.2018.11.001>[Get rights and content](#)

[Highlights](#)

* ?

We extract identity-disentangled representations for facial expression recognition (FER) without requiring expressive-neutral pairs in a testing task.

* ?

Proposing a novel recognition via generation scheme as a substitution of conventional hard sample mining.

* ?

The distance comparisons are largely reduced from $2K$ (triplet loss) to $2K$, where K is the number of sample in a training batch.

* ?

Our new architecture achieves the state-of-the-art on 3 popular FER datasets, which do not have neutral expression samples.

* ?

We alleviate the difficulty of threshold validation and anchor selection in conventional deep metric learning.

* ?

We learn distance metrics with much fewer distance calculations and training iterations, without sacrificing the performance.

* ?

We optimize the softmax loss and metric learning loss jointly based on their characteristics and tasks.

* ?

A novel approach to generate photorealistic and identity-preserved normalized face image.

Abstract

Various factors such as identity-specific attributes, pose, illumination and expression affect the appearance of face images. Disentangling the identity-specific factors is potentially beneficial for facial expression recognition (FER). Existing image-based FER systems either use hand-crafted or learned features to represent a single face image. In this paper, we propose a novel FER framework, named *identity-disentangled facial expression recognition machine* (IDFERM), in which we untangle the identity from a query sample by exploiting its difference from its references (e.g., its mined or generated frontal and neutral normalized faces). We demonstrate a possible recognition via generation scheme which consists of a novel hard negative generation (HNG) network and a generalized radial metric learning (RML) network. For FER, generated normalized faces are used as hard negative samples for metric learning. The difficulty of threshold validation and anchor selection are alleviated in RML and its distance comparisons are fewer than those of traditional deep metric learning methods. The expression representations of RML achieve superior performance on the CK + , MMI and Oulu-CASIA datasets, given a single query image for testing.

* Previous article in issue

* Next article in issue

Keywords

Hard negative generation

Adaptive metric learning

Face normalization

Facial expression recognition

1\ Introduction

Facial expression is the most natural and expressive nonverbal channel for humans to communicate their emotions [1]. Therefore, facial expression recognition (FER) has been an important and active topic for a wide range of applications including soft biometrics, digital entertainment, health care, robot systems and human-computer interaction (HCI). Ekman and Friesen postulated the universality of neutral (Ne) and six prototypical human facial expressions, namely, anger (An), disgust (Di), fear (Fe), happiness (Ha), sadness (Sa) and surprise (Su) [2].

The performances of the FER systems usually depend heavily on facial expression representations, which are affected by pose and illumination variations as well as facial morphology variations (i.e., identity-specific factors). As some facial expressions involve subtle facial muscle movements, the extracted expression-related information from different classes (in this paper, class refers to expression) can be overwhelmed by high-contrast identity-specific geometric or appearance features degrading FER performance. As illustrated schematically in Fig. 1, we want the intra-class distances of happy face images from different people to be smaller than the inter-class distances between face images of different expressions from the same subject. However, the nuisance identity factors often dominate the representation of the image in pixel space causing two images of the same subject with different expressions to be closer to each other than the same-expression images from two different subjects. This is because the extracted facial representation often contains identity-specific information that is irrelevant and may be counter-productive for the FER task. These identity-specific factors may degrade the FER performance on new identities unseen in the training data.

1. Download: [Download high-res image \(185KB\)](#)

2. Download: [Download full-size image](#)

Fig. 1. Illustration of the desired representations in the Face Expression Recognition (FER) feature space. The 'class' here refers to the facial expression.

Aided by the advances in deep learning for computer vision [3], much progress has been made on extracting a set of features to represent a single facial expression image [4], [5]. The hand-crafted features are constructed by exploiting domain knowledge of the specific relationships within pixels so that the features are invariant to some simple transformations (e.g., translation and scaling). More recently, feature-learning approaches are being investigated because of their ability to produce features that are tolerant to complex transformations. In the case of FER, identity-associated factors may fall into this category. Yet, unfortunately, due to the tight coupling of the various nuisance factors, when we try to reduce the sensitivity to identity with these state-of-the-art strategies, we are unlikely to satisfactorily disentangle all variations in facial morphology.

Our preliminary work [6], [7] proposed to improve the facial expression recognition performance by disentangling the identity factors in a face image through the metric learning method. To alleviate the slow convergence caused by trivial training samples in metric learning, [6] compared the query image to a negative set containing other facial expression images from the same subject as shown in Fig. 3. However, in practice, the structure of real-world FER datasets results in a significant constraint: the dataset may not contain images of every facial expression for every subject. Actually, we may not need to compare a query facial expression with any other class of facial expressions. According to some psychology and anatomy research [8], the muscle activities of different facial expressions initiate from the neutral face as illustrated schematically in Fig. 2. This is also the fundamental principle underlying the action units (AUs) and Facial Action Coding System (FACS) proposed by Ekman [2]. Since the expressive classes are naturally more discriminative from each other than from the neutral face, in the training stage, we may want to emphasize the neutral-expressive distance more than requiring large distance between those expressive classes. The neutral face images can be the ideal hard samples which can improve the learning efficiency of metric learning. However, expressive-neutral face image pairs of every person may not be always available in real world applications and in some FER

image datasets.

1. Download: Download high-res image (308KB)

2. Download: Download full-size image

Fig. 2. A schematic depicting the 6 basic facial expressions and their relationships to the neutral face. Representations in the outer ring correspond to higher-intensity facial expressions compared to those in the inner ring.

The above insights suggest that we should generate the frontal and neutral normalized face image of the query image and disentangle the possibly counter-productive identity information and we proposed _identity-disentangled facial expression recognition machine_ (IDFERM) towards this goal. IDFERM consists of two main parts, namely, the hard negative generation (HNG) network and the radial metric learning (RML) network. Specifically, given a query facial expression image, its normalized reference will be synthesized using an HNG network trained using expression-normalize face image pairs of the same subject. Then, the query-reference pairs are fed into the radial metric learning (RML) network, which uses an inception style convolutional (Conv) layers group and a unified two-branch fully connected (FC) layers framework to extract the contrast of query-reference pair by simultaneously optimizing the softmax loss and RML loss. By pushing the representation of facial expression images away from their generative references and pulling them close to their cluster centers of each expression, the RML can disentangle those nuisance factors to balance the intra- and inter- class variations.

Unlike other image-based FER systems that produce a representation from a single input image, RML utilizes the expression-reference pair to disentangle the identity-specific factors. In contrast to video-based FER methods [9], [10] or real facial expression image pair-based methods [6], [7], [11], our method employs synthetically generated references to address the real-world limitation that sometimes the dataset does not contain all possible facial expression examples for some subjects.

The preliminary versions of the concepts in this paper were published in the 2017 Biometrics workshop of _IEEE Conference on Computer Vision and Pattern Recognition_ [6] and 2018 IEEE International Conference on _Identity, Security and Behavior Analysis_ [7]. In this paper, we extend those basic concepts in the following ways:

* 1)

We investigate the prior relationship of different expression classes and propose a novel recognition via generation scheme as a possible substitution of conventional hard sample mining.

* 2)

We design an end-to-end IDFERM to extract identity-disentangled representations for FER without requiring real expression-neutral pair inputs in the FER datasets.

* 3)

The number of needed distance comparisons for adapted RML is orders of magnitude smaller than the number needed for conventional metric learning approach.

* 4)

We conduct all experiments using the new architecture and test on more datasets without neutral expression samples.

In summary, this paper makes the following contributions.

* ?

We propose a generalized metric learning loss function with adaptively learned reference threshold which alleviates the difficulty of threshold validation

and anchor selection.

* ?

With the identity-aware HNG for hard negative generation, the adapted RML learns distance metrics with fewer input iterations and distance calculations, without sacrificing the performance for identity-invariant FER.

* ?

We jointly optimize the softmax loss and metric learning loss in a unified two-branch FC layer metric learning CNN framework based on their characteristics and tasks.

* ?

The proposed HNG network is a novel approach to generate photorealistic and identity-preserved normalized face image by combining prior knowledge from data distribution and domain knowledge of faces.

* ?

Using numerical experiments, we demonstrate that the proposed method achieves promising results on CK + , MMI and Oulu-CASIA data sets.

The rest of this paper is organized as follows. Section 2 briefly reviews related work in the literature. Section 3 introduces in detail the proposed normalized face generation with perceptron generative adversarial networks. Section 4 shows its application to FER with RML network. Section 5 reports the experimental results and ablation study of the auxiliary parts in HNG network. Finally, Section 6 provides our conclusions.

2\ Related work

Despite receiving considerable research attention, FER remains very challenging [12]. Research developments in deep learning, especially the success of convolutional neural networks (CNN), have made high-accuracy image classification possible in recent years. Deep learning-based FER methods have emerged starting with Bengio et al. [13] who described the use of carefully designed CNN to learn expression features from raw pixels. Despite its popularity, current softmax loss-based approach does not explicitly reward intra-class compactness and inter-class separation, and identity-related factors remain major obstacles for FER. Machine recognition usually is based on similarity metrics, but those metrics may be more sensitive to identity than expressions. To decouple these two types of similarity and exploit the appearance information, substantial efforts have been dedicated to extracting features by learning [14]. Given that the expressions are formed by relaxing or contracting some facial muscles that result in temporally deformed facial features, identity-disentangled representations for FER normally separate a face with expression into a main component neutral face that encodes identity cues and an action component that encodes motion cues (such as movements of eye brows, cheeks, lips, eyelids and nose) which are related to the AUs and FACS [2].

FER is certainly not unique among computer vision applications that have to cope with nuisance factors causing variability in the data. Deep metric learning approaches have been shown to be successful for person and vehicle identification tasks [15], [16], [17], which also exhibit large intra-class variations. The initial work in this domain [18] involves training a Siamese network. Pairwise examples are fed into two symmetric sub-networks and the network is updated using contrastive loss function, i.e., their extracted representations should be close to each other if the inputs have the same class label, otherwise the distance between these representations should be large. One improvement is the triplet loss [19], in which, the inputs are triplets, each consisting of a query, a positive example and a negative example. An anchor is chosen from the query or positive examples, then the

method requires the difference of the distance from the anchor point to the positive or query example and from the anchor point to the negative example to be larger than a fixed margin τ . Recently, some variants of this offering faster and more stable convergence have been developed. The $(N+1)$ -tuple loss [20] incorporated multiple negative examples while the coupled cluster loss (CCL) [15] incorporated multiple positive examples in a tuple. The center of positive examples is set as the anchor in CCL. By comparing each example with c^+ instead of each other, the number of distance evaluations needed are reduced significantly.

For the situation shown in Fig. 3, the triplet loss, $(N+1)$ -tuple loss and CCL are all 0, since the distances between the anchor and positive examples are indeed smaller than the distance between the anchor and negative examples for a margin τ . This means the loss function will learn to neglect such non-trivial samples. We will need many more input iterations with properly selected anchors to correct this situation. The fixed threshold in the contrastive loss was also proven to be sub-optimal [21]. The difficulty of threshold validation and anchor selection have long been significant challenges until our initial work [6], which included an adaptive $(N+M)$ -tuple clusters loss function.

1. Download: [Download high-res image \(249KB\)](#)

2. Download: [Download full-size image](#)

Fig. 3. Failed case of (a) triplet loss, (b) $(N+1)$ -tuple loss, and (c) Coupled clusters loss. The preliminary $(N+M)$ -tuple clusters loss (d) uses two thresholds to avoid the anchor selection issue and threshold-parameters τ_i , τ_t and τ , do not need manual tuning [6], [7]. We use x^+ (yellow points) and x^- (squares) to denote the positive and negative examples of a query example x , meaning that x^+ is the same class of x , while x^- is not. $f(\cdot)$ is an embedding kernel.

Also, the traditional online or offline mini-batch sample selection is a large additional computational burden and can result in poor local optima [22]. Generating all possible pairs or triplets would result in quadratic and cubic complexity, respectively and most of these pairs or triplets are not very useful for the training [6]. Our initial work [6] utilized identity-aware hard-negative mining and online positive mining for FER, but it still suffers from the dataset-sensitive and computationally-expensive example mining to provide nontrivial tuples.

Several approaches have been proposed for generative models. Conventional methods such as Principal Components Analysis (PCA), Independent Components Analysis (ICA), Gaussian Mixture Model (GMM), etc., have difficulty in modeling complex patterns of irregular distributions [23]. Recently, Restricted Boltzmann machines (RBM), Hidden Markov Model (HMM), Markov Random Field (MRF) etc., have been employed for modeling images of digits, texture patches, and well-aligned faces [24]. However, the limited ability of feature representations restricts further development. Since deep hierarchical architectures of the recent generative models are capable of capturing complex structure of data, generated images from these deep hierarchical structures are more realistic. The denoising auto-encoder (DAE) pairs a differentiable encoder and decoder, which encodes an image sample x to a latent representation z and then decodes the z back to another image x' [25]. For the normalized face generation task, pose and expression are regarded as the noise to be mitigated. The main limitation of this approach is that the squared pixel-wise reconstruction error would cause the generated samples to look blurry as they generate the mean image of the distribution. Generative Adversarial Network (GAN) [26] simultaneously trains two networks: a

generative network $_Gen_$ to synthesize images (maps latents $_z_$ to image space), and a discriminative network $_Dis_$ to discriminate between real training images from generated images. With the GAN, an expected image can be generated from a randomly sampled vector $_z_$ from a certain distribution. Normally, the GAN schemes are not well-matched to supervised recognition tasks. The GAN-generated results are expected to align with the central part of the data distribution, while the boundary between classes in feature space is more important for classification. Limited research has been devoted to this topic. The semi-GAN [27] adds an extra task for a discriminator network to improve semi-supervised recognition task. The face rotator schemes proposed by Tran [28] generates a frontal face as the preprocessing for the face recognition network.

3\ Hard negative generation

As we are trying to disentangle identity-related factors from a facial expression image $_x_$, a reference neutral face image from the same subject is required to obtain the difference between the neutral face and $_x_$ for FER. However, such a neutral face reference image is not always available in real world application scenarios. Instead of mining several negative samples, we directly use the generated normalized face image as negative sample. The goal of the hard negative generation (HNG) network is to produce a photorealistic and identity-preserved normalized face image $x?$ from the probe image $_x_$. The network architecture and loss functions are illustrated in Fig. 4.

1. Download: [Download high-res image \(554KB\)](#)

2. Download: [Download full-size image](#)

Fig. 4. Framework of our IDFERM, in which, the left and right side are the HNG and RML network, respectively. We feed the input image $_x_$ to an encoder-decoder structure and generate its transformed output $x?$. The $_Dis_$ is trained to determine if its input is from the guiding set (real image) or the encoder-decoder structure (generated image), thus encouraging the encoder-decoder structure to generate images more similar to the images in our guiding set. The guiding set contains all of the target images $_y_$, which is a compilation of numerous frontal and neutral real face images. The Light-CNN is used to extract the identity feature for identity similarity measurements, and the VGG-Face is adopted to embed image for feature level perceptual similarity measurements. The mined real expressive face images and their corresponding generated normalized face images are feeded to the RML, an adaptive deep metric learning framework, to disentangle the identity factors in a face image for FER by minimizing both the softmax (cross-entropy) loss and RML loss. Our HNG network is composed of five major components: 1) an Encoder network $_Enc_$, (2) a Decoder network $_Dec_$, (3) a Discriminator network $_Dis_$, (4) the $_Light\ CNN_$ network and (5) the $_VGG-facenet_$. The function of $_Enc_$ and $_Dec_$ network is the same as that in denoising auto-encoder [25]. In DAE, the output is not required to be exactly the same as the input. For example, the denoising auto-encoder takes in an image that has been corrupted by some form of ?noise?, and is forced to output a denoised version of that image by requiring the output image to be similar to the original ?clean? image. In our application, a face image with expression (input image $_x_$) can be regarded as a copy of neutral face image (target image y) that has been corrupted by expressions. The denoising (disentangling expression from a face image) is achieved by requiring our output $x?$ to be close to the target image y , as the DAE requires its generated image to be close to a clean target image instead of the original noisy image. The $_Enc_$ maps the input sample image $_x_$ to a latent representation $_z_$ through a learned distribution $P(_z|x_)$, while the $_Dec_$ generates predicted a facial image $x?$ corresponding to $_z_$. The function

of the $_Dec_$ and $_Dis_$ is the same as that in the GAN [26]. The $_Dec_$ network tries to generate the real distribution by the loss of $_Dis_$ which learns to distinguish between generated image x and real image in the guiding set $_g_$. The guiding set contains all of the target images $_y_$, which is a compilation of numerous frontal and neutral real face images. The input-target pairs $\{x_i, y_i\}$ from multiple identities are required to learn the parameters $_\theta_$ of the differentiable encoder $_? Enc_$ and decoder $_? Dec_$, where $_x_$ is a face image with expression and $_y_$ is the frontal neutral face image of that person. In our experimental setting, five different loss functions are used to combine the advantages of high quality GAN and stable auto-encoder which encodes the data into a latent space $_z_$. In this section, we show how the multiple objective functions are employed for different parts to generate the facial reference images for FER.

3.1. Feature space perceptual loss

The squared error loss between the CNN feature representations is adopted to represent the feature-level perceptual loss. We make use of the independently-trained and fixed VGG-FaceNet [29] to model this semantic feature-level loss. Although it comprises 14 Conv layers and 3 FC layers, we omitted the deeper layers after the 5th Conv layer because their limited spatial resolution cannot support good image reconstruction performance. Denoted by $_l_$, the feature map of the $_l_$ th convolution layer of VGG-FaceNet is used to extract the feature representations using the standard forward-propagation process. The semantic perceptual loss between two images x and y on the $_l_$ th convolutional layer is defined as the following squared-error loss between the two feature maps.

$$L_{feat}(x, y) = \frac{1}{W \times H} \sum_{n=1}^W \sum_{m=1}^H \|I_{l,n,m}(x) - I_{l,n,m}(y)\|^2$$

where $_W_$ and $_H_$ denote the width, height of the $_l_$ th feature map and $_I_{l,n,m}_$ is the value of the $_l_$ th feature map at point (n, m) . In our experiment, $_l_ = 5$ based on empirical testing.

3.2. Symmetry loss

Symmetry is an inherent property of normal human faces. The symmetry loss of a face image takes the form:

$$L_{sym}(x) = \frac{1}{H \times W} \sum_{n=1}^W \sum_{m=1}^H |x(n, m) - x(W - n + 1, m)|$$

where $_W_$ and $_H_$ are the width and the height of the images and (n, m) denotes the pixel of the generated image. The $| \cdot |$ denotes the absolute value. For simplicity, when training our model with the symmetry loss, all the inputs are aligned and detected if the occluded parts are on the right side of image. If not, images are flipped so that the occluded parts are on the right side. Real-world images may not exhibit the strict symmetry of pixel values. Considering the consistency of the pixel difference inside a local area, and the gradients at a point along all directions are largely preserved under different illuminations, minimizing a symmetry loss in the Laplacian space should emphasize human faces.

3.3. Adversarial loss

We introduce a discriminator $_Dis_$ which serves as a supervisor to push the synthesized image to reside in the manifold of frontal neutral face images. It can reduce the blur effect and produce visually pleasing results.

The $_Dis_$ aims to discriminate the predicted frontal neutral face image x_i from real ones $_g_i_$ in the guiding set, and is trained concurrently with the transform network ($_Enc_$ and $_Dec_$). The transform network tries to "trick" the $_Dis_$ to classify the generated images as real. Formally, the discriminator is trained to minimize the following binary cross entropy loss:

$$L_{GAN}(Dis(g_i, x_i)) = -\log(Dis(g_i)) - \log(1 - Dis(x_i))$$

With respect to $_Dec_$, the parameters are trained by minimizing the following

loss:(5) $LGAN_{Dec}(x_j) = -\log(Dis(x_j))$

3.4. Identity-Preserving loss

Synthesizing the frontal neutral face image while preserving the identity is a critical part of IDFERM. We introduce a direct supervision to reward the perceptual similarity between input and generated images using the face verification network. In our approach, we use the pre-trained Light CNN, a compact network that has only 4 convolution layers with Max-Feature-Map operations and 4 max-pooling layers [30]. In this work, the identity-preserving loss is defined based on the activations of the last two layers of the Light CNN: $L_{id} = \frac{1}{2} \sum_{l=121}^L \sum_{n=1}^W \sum_{m=1}^H |l, n, m(x) - l, n, m(y)|$ where l , H , l denotes the width and height of the l -th layer, n , m is the value of the feature map (n, m) point and $|x|$ denotes the absolute value.

3.5. Pixel-wise loss

Adversarial training is known to be sensitive to hyper parameters. Adding the following pixel-wise L1 loss $L_{pixel} = \sum_{n=1}^W \sum_{m=1}^H |x_{n,m} - y_{n,m}|$ in the image space with a relatively small weight is one method to stabilize the training and accelerate the optimization. $x_{n,m}$ and $y_{n,m}$ are the pixel level gray values of the (n, m) -th pixel.

Using judicious selection of aforementioned loss functions, we train the Enc , Dec and Dis simultaneously. The error signal from adversarial loss and symmetry loss are not back-propagated to Enc . Several tradeoff parameters constrained between 0 and 1 are used to balance the aforementioned loss functions. The weights α_1 and α_2 shown in Algorithm 1 are the tradeoff parameters for L_{feat} , L_{id} and L_{pixel} for the Enc and Dec . The parameter α_3 is used to weight the L_{sym} in Dec . As Dec also receives the error signal from the Dis , a parameter α_4 is used to weight the ability of fooling the discriminator.

Algorithm 1. Training the HNG network.

```
**  $\alpha_1$   $\alpha_2$   $Enc$ ,  $\alpha_3$   $\alpha_4$   $Dec$ ,  $\alpha_5$   $\alpha_6$   $Dis$  *** initialize network parameters
---
**Repeat**
**  $X$  ** ? random mini-batch from dataset
**  $Z$  ** ?  $Enc(X)$ 
 $X$  ? ?  $Dec(Z)$ 
 $L_{feat} = \alpha_1 \sum_{l=121}^L \sum_{n=1}^W \sum_{m=1}^H |l, n, m(x) - l, n, m(y)|$ 
 $L_{sym} = \alpha_2 \sum_{n=1}^W \sum_{m=1}^H |x_{n,m} - y_{n,m}|$ 
 $LGAN_{Dis} = -\log(Dis(g_i)) - \log(1 - Dis(x_j))$ 
 $LGAN_{Dec} = -\log(Dis(x_i))$ 
 $L_{id} = \frac{1}{2} \sum_{l=121}^L \sum_{n=1}^W \sum_{m=1}^H |l, n, m(x) - l, n, m(y)|$ 
 $L_{pixel} = \sum_{n=1}^W \sum_{m=1}^H |x_{n,m} - y_{n,m}|$ 
**/** Update parameters according to gradients
**  $\alpha_1$   $\alpha_2$   $Enc$  ????  $Enc(L_{feat} + \alpha_1 L_{id} + \alpha_2 L_{pixel})$ 
**  $\alpha_3$   $\alpha_4$   $Dec$  ????  $Dec(L_{feat} + \alpha_1 L_{id} + \alpha_2 L_{pixel} + \alpha_3 L_{sym} + \alpha_4 LGAN_{Dec})$ 
**  $\alpha_5$   $\alpha_6$   $Dis$  ????  $Dis(LGAN_{Dis})$ 
**Until** deadline
```

We show some of the input-output pairs of our HNG network in Fig. 5. As the common quantitative metrics (e.g., log-likelihood of a set of validation samples) are often not very informative for perceptual generative models [31], we provide a qualitative comparison of visual quality and a quantitative evaluation of identity-preservation in Section 5.

1. Download: [Download high-res image \(466KB\)](#)

2. Download: [Download full-size image](#)

Fig. 5. Input-output pairs of the proposed reference generation (HNG) network

from the CK + , MMI and Oulu-CASIA dataset. Top row in each pair: a subject in a database with different expressions (from left to right: angry, disgust, fear, happy, neutral, sad and surprise). Bottom row in each pair: generated normalized face images from the input of the expressive face images in the row above.

Unlike previous generative methods that utilize their intermediate features for the recognition tasks, the resulting expression- and pose- disentangled face image has potential for several downstream applications, such as facial expression or face recognition, and attribute estimation.

4\ Radial metric learning

The proposed RML only requires the comparison of the representation of the query sample f_i with the representation of its generated reference $f_i^?$ and its cluster center C_{yi} . We introduce a distance T from the query sample x_i to control the relative boundary $(T+2)$ and $(T-2)$ for the intra-class center and generated references, respectively. The RML loss function is formulated as follows. (8) $L(\{x_i\}_{i=1}^K, \{x_i^?\}_{i=1}^K; f) = 1/K \sum_{i=1}^K \{ \max(0, D(f_i, C_{yi}) - T + 2) + \max(0, T + 2 - D(f_i^?, C_{yi})) \}$ Only if the distances from all online mined examples f_i to its updated C_{yi} are smaller than $(T-2)$ and the distances from all the generated references $f_i^?$ to its updated C_{yi} are larger than $(T+2)$, the loss $L(\{x_i\}_{i=1}^K, \{x_i^?\}_{i=1}^K; f)$ can get a zero value. A simplified geometric interpretation of this is shown in Fig. 6.

1. Download: Download high-res image (300KB)

2. Download: Download full-size image

Fig. 6. The proposed radial metric learning (RML) framework. A small circle without border is the representation of a sample (i.e., facial expression image) and the different classes are represented by different colors. The small gray circles with colored border are their corresponding generated references (i.e., normalized face images). The orange points with colored border are the cluster centers of each classes. The big dashed circles are the boundaries of each classes in the feature space, which are expected to have small radius and far away from each other.

By assigning different values for T and γ , we define a flexible learning task with adjustable difficulty for the network. We do not use the special case that requires inter-class variation to be zero (i.e., $T = \gamma/2$) as the center loss [32] for the FER training set usually contains some unreliable labels [33]. However, these two hyper-parameters need manual tuning and validation. In here, we formulate the reference distance T to be a function $S(\cdot, \cdot)$ which should be trained automatically, instead of a constant.

Inspired by the Mahalanobis distance matrix M in Mahalanobis distance D (Eq. 9), which is a positive semi-definite (PSD) matrix and can be calculated via the linear fully connected layer as in [34], we try to automatically train both S and D . Since the difference of the reference distance and the distance function need to be calculated in two terms in Eq. 8, a possible solution is to calculate $(S-D)$ function via the linear FC

layer. (9) $D(f_1, f_2) = \sqrt{(f_1 - f_2)^T M (f_1 - f_2)}$

Since the metric M itself is quadratic, we assume that S has a simple quadratic form: (10) $S(f_1, f_2) = 1/2 f_1^T A f_1 + 1/2 f_2^T A f_2 + f_1^T B f_2 + c(f_1, f_2) + b$ where A and B are both the $d \times d$ real symmetric matrices (not necessarily positive semi-definite), c is a d -dimensional vector, and b is the bias term.

Then, a new quadratic expression $H(f_1, f_2) = S(f_1, f_2) - D(f_1, f_2)$ is defined to combine the reference distance function S and the Mahalanobis distance metric function D . Substituting $S(f_1, f_2)$ and $D(f_1, f_2)$ into $H(f_1, f_2)$, we

get: (11) $H(f_1, f_2) = 1/2 f_1^T (A - 2M) f_1 + 1/2 f_2^T (A - 2M) f_2 + f_1^T (B + 2M) f_2 + c(f_1, f_2) + b$ (12) $H(f_1, f_2) = 1/2 f_1^T A f_1 + 1/2 f_2^T A f_2 + f_1^T B f_2 + c(f_1, f_2) + b$

$$1tBf2+ct(f1?f2)+b\text{where}$$

$A=(A^{?2}M)$ and $B=(B^{?2}M)$. Suppose A is PSD and B is negative semi-definite (NSD), A and B can be factorized as $LATLA$ and $LBTLB$. Then $H(\underline{f}_1, \underline{f}_2)$ can be rewritten as

$$\text{follows:} (13) H(f_1, f_2) = 12f_1tLAf1 + 12f_2tLAf2 + f_1tLBtLBf2 + ct(f1?f2) + b \quad (14) H(f_1, f_2) = 12(LAf1)t(LAf1) + 12(LAf2)t(LAf2) + (LBf1)t(LBf2) + ct(f1?f2) + b$$

Motivated by the above, we propose a general, computationally feasible loss function. Following the notations in the preliminaries and denoting $(LA, LB, \underline{c})t$ as W which can be learned via the linear fully connected layer, we have: (15) $L(W, \{x_i\}_{i=1}^K, \{x_i^?\}_{i=1}^K; f) = 1K \sum_{i=1}^K \{ \max(0, \sum_{j=1}^K H(f_i, C_{yi})) + \max(0, H(f_i^?, C_{yi}) + \sum_{j=1}^K \dots) \}$

Moreover, we simplify $\sum_{j=1}^K$ to be the constant 1, since changing it to any other positive value results only in the matrices being multiplied by a corresponding factor. Our hinge-loss like function is given as follows. (16) $L(W, \{x_i\}_{i=1}^K, \{x_i^?\}_{i=1}^K; f) = 1K \sum_{i=1}^K \{ \max(0, 1 \sum_{j=1}^K H(f_i, C_{yi})) + \max(0, H(f_i^?, C_{yi}) + 1) \}$

By doing this, the adaptive threshold can be seamlessly factorized into a linear-fully connected layer for end-to-end learning [34]. The RML loss can also be easily used as a drop-in replacement for the triplet loss and its variants, as well as used in tandem with other performance-boosting approaches and modules, including modified network architectures, pooling functions, data augmentations or activation functions.

For a training batch consisting of \underline{K} query samples, the number of input passes required to evaluate the necessary embedding feature vectors in our application is \underline{K} , and the total number of distance comparisons can be $2 \times \underline{K}$. Normally, \underline{K} is much larger than 2. In contrast, triplet loss and $(\underline{N} \setminus + 1)$ -tuple loss require $\underline{O}(\underline{K} \times 3)$ comparisons, the contrast loss and CCL require $\underline{O}(\underline{K} \times 2)$ comparisons, and the $(\underline{N} + \underline{M})$ -tuple cluster loss requires $2(\underline{N} + \underline{M}) \times \underline{K}$ comparisons after a strict example mining scheme using the special structure of some FER datasets (i.e., each subject has all 6 expressions). Here \underline{N} and \underline{M} are the number of mined positive samples and the number of mined negative samples, respectively. Even for a dataset of a moderate size, it is computationally impractical to load all possible meaningful triplets into the processor memory for model training. With predefined anchors (i.e., C_{yi} and $f_i^?$), we also alleviate the difficulty of anchor selection [6].

The inception convolutional FER network and two-branch FC layer joint metric learning architecture proposed in our preliminary paper [6] are used in our framework in Fig. 4. The convolutional groups are made up of a 1×1 , 3×3 and 5×5 Conv layers in parallel.

Combining the metric learning loss and softmax loss is an intuitive idea to possibly achieve better performance [35]. However, combining them directly on the last FC layer is sub-optimal. The basic idea of building two-branch FC layers after the deep convolution groups is to combine two losses at different levels of tasks. We learn the detailed features shared between the same expression class with the expression classification (EC) branch, while exploiting semantic representations via the metric learning (ML) branch to handle the significant appearance changes from different subjects. The connecting layer embeds the information learned from the expression label-based detail task to the identity label-based semantic task, and balances the weights in the two task streams. This type of combination can effectively alleviate the interference of identity-specific attributes. The inputs of connecting layer are the output vectors of the former FC layers- \underline{FC}_2 and \underline{FC}_3 , which have the same dimension denoted as \underline{D} input. The output of the connecting layer, denoted as \underline{FC}_4 with dimension \underline{D} output, is the feature vector fed into the second layer of the ML branch. The connecting

layer concatenates two input feature vectors into a larger vector and maps it into a $2D_{out}$ output dimension space: $(17)FC2 \times 4 = P1[FC2 \times 2; FC2 \times 3] = P1tFC2 \times 2 + P2tFC2 \times 3$ where P is a $(2D_{in} \times 2D_{out})$ matrix, $P1$ and $P2$ are $D_{in} \times D_{out}$ matrices.

Regarding the sampling strategy, every training image is used as a query example in an epoch. In practice, the softmax loss will only be calculated for the query examples. The relative importance of the two loss functions is managed by a weight λ . During the testing stage, this framework takes one query image and its generated reference image as input, and determines the classification result through the EC branch with the softmax loss function. Our disentangled feature learning scheme is described in Algorithm 2.

Algorithm 2. Disentangled feature learning algorithm.

Input

Randomly chose K query examples $\{x_i\}_{i=1}^K$

and their generated references $\{x_i^*\}_{i=1}^K$

Output: The parameters of the FER network θ_{FER}

1. while not converge do

2. map examples to feature plane with CNN to get: $\{f_i\}_{i=1}^K$ and $\{f_i^*\}_{i=1}^K$

3. calculate the cluster centers C_i for each class

4. LRML $\lambda \sum_{i=1}^K \{\max(0, H(f_i, C_i)) + 1 + \max(0, H(f_i^*, C_i)) + 1\}$

5. $L_{softmax} = -\sum_{i=1}^K \log(\frac{e^{f_i^T C_i}}{\sum_{j=1}^K e^{f_j^T C_j}})$

6. Compute the joint loss $L_{softmax} + \lambda LRML$

7. Compute the backpropagation error

8. Update the parameters

End while

5\ Numerical experiments

In this section, we compare the IDFERM with state-of-art methods on three benchmark datasets, i.e., CK + , MMI and Oulu-CASIA datasets. Details of our training data are provided in Section 5.1, followed by our preprocessing methods in Section 5.2, and the implementation details in Section 5.3. In Section 5.4, we report a series of ablation experiments to analyze the function of our auxiliary networks. Numerical experiment results are shown in Section 5.5.

5.1. Training data

Besides the FER datasets, a variety of large datasets of facial images for face recognition are publicly available online. We give some samples from the VGG-Face dataset Fig. 7. We use the VGG-Face dataset [29] to extend our data for neutral face generation. It contains approximately 2.6 million face images, but very few of these fit our requirements of neutral expression, front-facing, having no occlusion, and of sufficient resolution for face region. We use the Google Cloud Vision API to remove those images that look blurry, with high emotion score or eyeglasses or tilt or pan angles beyond 5° . These frontal and neutral face images are used as our target and guiding set samples. Their corresponding non-compliant images from the same subject are used as the inputs. All the samples are aligned and cropped to 64×64 Gy images. After filtering, we have about 12 K target images ($< 0.5\%$ of the original set) and 50 K input-target pairs. These data are used for pretraining to initialize the network parameters and then fine-tuned using the CMU Multi-PIE [36].

1. Download: Download high-res image (541KB)

2. Download: Download full-size image

Fig. 7. Samples from the VGG-Face dataset. Each row contains the face images

of the same person.

The CMU Multi-PIE itself is a facial expression dataset, but the facial expression labels it uses are slightly different from modern expression classification system. There are 4 expressions are useable (114,305 neutral images, 19,817 surprise images, 22,696 disgust images and 47,388 happy images), while the squint and scream are not regarded as expression now. Some samples are shown in Fig. 8. It contains more than 750,000 images of 337 people taken from fifteen directions, and in nineteen illumination conditions. There are four recording sessions in which subjects were instructed to display neutral, happy, disgust and surprise facial expressions. It is more close to our FER dataset in testing stage than the filtered VGG face dataset. We selected only the five groups of the nearly frontal view faces (-45° to $+45^\circ$). The neutral images from the 0° view are used as our target image and the guiding set.

1. Download: Download high-res image (195KB)

2. Download: Download full-size image

Fig. 8. Samples from the CMU Multi-PIE dataset. Each row contains images of the same person.

5.2. Preprocessing

We follow the [6] to locate the 49 facial landmarks. Then, face alignment is done to reduce in-plane rotation and crop the region of interest based on the coordinates of these landmarks to a size of 64×64 . An augmentation procedure is employed to increase the number of training images and alleviate the chance of over-fitting. We crop five 60×60 size patches from the center and four corners, flip them horizontally and transfer them to grayscale images. All the images are processed with the standard histogram equalization and linear plane fitting to remove unbalanced illumination. Finally, we normalize them to have zero mean and unit variance. In the testing phase, a single center crop with the size of 60×60 is used as input data.

5.3. Implementation details

We use 64×64 Gy images as the input-target pairs for the neutral face generation training. The filtered VGG-FaceNet and Multi-PIE images are used to pre-train the neutral face generation network. We construct the guiding set using the filtered VGG-FaceNet frontal neutral view and the 0° view neutral images from the Multi-PIE. Following the experimental protocol in [6], we pre-train our inception style convolutional groups, two branch FC layers on with 204,156 frontal view (-45° to 45°) face images selected from the CMU Multi-PIE dataset for 300 epochs, optimizing the joint loss using stochastic gradient descent with a momentum coefficient of 0.9. The initial network learning rate, batch size, and weight decay parameter are set to 0.1, 128, 0.0001, respectively based on optimizing the parameter choices using the validation set. If the training loss increased by more than 25% or the validation accuracy does not improve for ten epochs, the learning rate is halved and the previous network with the best loss is reloaded. We select the highest accuracy training epoch as our pre-trained model. In the fine-tuning stage, the mini-batch set size is fixed to two times the number of expression classes of the dataset. Random search is employed to select 2 images from each expression class to form the mini-batch set. The tuple-size is set to 12. In all our experiments, we set $\lambda = 0.1$, $\lambda_1 = 3 \times 10^{-2}$, $\lambda_2 = 10^{-2}$, $\lambda_3 = 0.3$ determined by manual tuning. The weight of joint learning $\lambda = 1$. In the testing phase, only the convolutional groups and expression classification branch with softmax are used to recognize a single facial expression image.

The details of the encoder of HNG can be found in [7]. We fixed the latent vector dimension to be 256 and found this configuration to be sufficient for

generating images for FER. A series of fractional-stride convolutions (FConv) transforms the 256-dim vector $z \in \mathbb{R}^{256}$ into a synthetic image $x \in \mathbb{R}^{64 \times 64}$, which is of the same size as x_{real} . To further incorporate the prior knowledge of the frontal neutral face's distribution into the training process, we introduce a discriminator $_{\text{Dis}}$ to distinguish the generated face image from the real images in the guiding set.

The Leaky ReLU nonlinearities [37] are used in some Conv layers, where $\text{LReLU}(x) = \max(x, 0) + \lambda \min(x, 0)$. In our experiments, we set $\lambda = 0.1$. Optimizing this minimax objective function will continuously push the output of the generator to match the target distribution of the guiding set thus making the synthesized facial images to be more photorealistic. All the CNN architectures are implemented with the widely used deep learning tool TensorFlow [38].

5.4. Ablation study

The Light CNN and the first five layers of the VGG-FaceNet are used to embed the input, target or output images for the similarity measurements in different feature spaces. It is obvious that these two networks incur additional computation cost. We show in this section that they are needed. The difference of our models trained with and without the L_{id} is subtle in visual appearance, as can be seen in Fig. 9, but its effect on improving the identity likeness of the generated faces can be measured by evaluating the similarity of the input-outputs pairs using VGG-FaceNet. Fig. 10 shows the distributions of L2 distances between the embeddings of the facial expression images and their corresponding synthesized results, for models trained with and without this loss. Schroff et al. [18] consider two FaceNet embeddings to encode the same person if their L2 distance is less than 1.242. All of the synthesized images using the identity-preserving loss pass this test using FaceNet, but about 2% of the images would be identified as a different subject by FaceNet when not using the identity-preserving loss. We investigated the effect of the weight of identity preserving loss and show the identity inconsistent percentage in Table 1.

- 1. Download: Download high-res image (769KB)
- 2. Download: Download full-size image

Fig. 9. Examples of input, target, generated normalized face by RG network, generated normalized face by RG network without identity-preserving loss and reconstructed neutral face using only the auto-encoder (AE) structure. Real images are from CK + dataset which has an additional class called contempt (Co) class.

- 1. Download: Download high-res image (124KB)
- 2. Download: Download full-size image

Fig. 10. Histogram of VGG-Face net L2 error between the input face and the normalized pairs on the FER data collection. Blue: with the identity preserving loss which calculated by the Light CNN. Orange: without the identity preserving loss. The 1.242 threshold was used by Schroff et al. [18] to cluster identities in the LFW dataset. Without the Light CNN, about 2% of the generated neutral faces would not be considered from the same subject as the query faces.

Table 1. Percentage of identity errors as a function of the λ , the weight of identity-preserving loss term.

λ	0	0.001	0.005	0.01	0.015	0.02	0.025	0.03	0.035	0.04	0.05
Percentage (%)	2.13	1.77	1.06	0.32	0.13	0.11	0.1	0.1	0.1	0.1	0.1

The VGG-FaceNet is employed to calculate the feature level perceptual loss,

which is expected to make the generated result to keep more perceptually important image attributes, for example sharp edges and textures. This loss

was empirically given the largest weight in our experiments. In practice, without this part, we could never avoid the collapse of the adversarial training to generate the human face structure. We also analyzed the effect of different hyper-parameter values in RML. The parameter λ is used to balance the softmax loss and metric learning loss. We can see from the Fig. 11 (reproduced below) that the highest accuracy is achieved when $\lambda \in [0.95, 1]$. As can be seen in Fig. 12 below, the networks were not sensitive to $\lambda \in [0.075, 0.125]$.

1. Download: Download high-res image (181KB)
2. Download: Download full-size image

Fig. 11. Facial recognition accuracy on CK + (7-class) dataset as a function of λ , the parameter used to balance the softmax loss and metric learning loss.

1. Download: Download high-res image (174KB)
2. Download: Download full-size image

Fig. 12. Facial recognition accuracy on CK + (7-class) dataset as a function of parameter λ in LReLU.

5.5. Experimental results of FER

To evaluate the effectiveness of the proposed method, extensive experiments have been conducted on four well-known publicly available facial expressions datasets: CK + , MMI and Oulu-CASIA. Resulting IDFERM confusion matrices are shown in Fig. 13.

1. Download: Download high-res image (727KB)
2. Download: Download full-size image

Fig. 13. Confusion matrix of the proposed IDFERM evaluated in the (a) CK + seven-class, (b) CK + eight-class, (c) MMI and (d) Oulu-CASIA database. The predicted labels and the ground truth labels are shown in the ordinate and abscissa, respectively.

CK + Dataset [39]: We conducted both seven-class and eight-class expression recognition experiments (i.e., without or with neutral expression). In the setting without neutral sample, we directly compare to the most nontrivial hard negative samples (i.e., generated normalized face), which not only relaxes the requirements on the dataset (i.e., needing images of all different expressions of the same person) to extract the identity-disentangled expression representation but also reduces the number of comparisons in training stage. The training time of metric learning part is largely reduced as shown in Table 2. We note that [7] and IDFERM do need an additional training stage (around 6 hours) for normalized face generation, but the trained HNG network works for all of the FER datasets in our experiments without fine-tuning and can be regarded as a ready-made tool for several downstream tasks.

Table 2. Comparison of the metric learning training time on the datasets with a Titan X GPU.

Metric Learning Training Time	CK + (7-class)	CK + (8-class)	MMI	Oulu-CASIA
-------------------------------	----------------	----------------	-----	------------

---	---	---	---	---

Triplet Loss	5 h 24mins	?	3 h 14mins	?
2B(N + M) Softmax [6]	1 h 47mins	?	54 mins	?
Data Augmentation [7]	3 h 11mins	?	2 h 8mins	?
IDFERM	42 mins	1 h 4 mins	30 mins	56 mins

In the testing stage, IDFERM recognizes a query facial expression image in about 50 ms, which is satisfactory for many applications. Video-based methods normally need a relatively longer sampling time (>0.25 s) to collect the whole

expression change session.

From Table 3, we can see that the identity-disentangled representation with adaptive metric learning methods achieve higher accuracy than previous works. Therefore, comparing the query sample with its generated normalized face rather than all its other expressive query faces as in [6] is more efficient, which is consistent with the relationship of those expressions as analyzed in Fig. 2. However, just adding the generated neutral face images to original dataset enlarged the number of comparisons [7]. As an efficient hard negative mining scheme, the HNG offers the most nontrivial hard negative samples and the RML can efficiently utilizes them and outperforms the other methods. Table 3. Performance compares of the rank-1 recognition accuracy on the CK + dataset in terms of 7 expressions and the MMI dataset (without neutral expression).

Methods	CK + (seven-class)	MMI
MSR [13]	91.4%	N/A
BNBN [42]	96.7%	N/A
IBCNN [4]	95.1%	N/A
STM [43]	96.3%	N/A
CER(video) [44]	92.34%	70.12%
CDMML(video) [45]	96.6%	N/A
STMExplet(video) [9]	94.19%	75.12%
DTAGN(video) [10]	97.25%	70.2%
IACNN [11]	95.37%	71.55%
2B(N + M) Softmax [6]	97.1%	78.53%
Data Augmentation [7]	97.49%	80.26%
IDFERM	98.35%	81.13%

The improved accuracy compared to the other methods is appealing in the image-based 7-class CK + , MMI and Oulu-CASIA setting which do not have real-neutral samples as training data. It also generalized well in dataset with neutral expressions as shown in Table 4. With the added generated samples, the accuracy of neutral class is improved to 99% as shown in Fig. 13(b).

Table 4. Performance comparison of the rank-1 recognition accuracy on the CK + dataset in terms of the 8 expressions (with neutral expression).

Methods	CK + (eight-class)
AUDB [46]	93.70%
CNN + AD [47]	96.4%
FN2EN [5]	96.8%
IDFERM	97.76%

Benefitting from the generated data, the proposed IDFERM outperforms our earlier approach [6] by 2.6% on MMI dataset and 1.25% on CK + dataset. Using the same generated images as in [7], we achieve 0.87% and 0.86% improvements on MMI and CK + datasets respectively, and the number of comparisons per training batch is reduced from $2(_N + _M)_K$ (the $_N$ and $_M$ in MMI are 5 and 5 respectively [7]) to $2K$. As a consequence, the training time of the earlier approach [7] for MMI dataset is 2hours 8mins, while the IDFERM needs only 30mins on a single Titan X GPU as shown in Table 2. Considering the improved performance and reduced training time, the IDFERM is significantly more efficient than [7] to utilize the generated data.

Limited training data has long been a challenge for facial expression recognition. For example, [11] utilized the FER-2013 dataset for pre-training and then fine-tuned their facial expression recognition network in CK + /MMI datasets. The FER-2013 dataset is even larger than the Multi-PIE dataset. We

chose the Multi-PIE pretraining for fair comparison with previous works [6]. We added a comparison of recognition accuracy with/without the pre-training, and show the results in Table 5. We can see that the pre-training can improve the performance consistently.

Table 5. Comparison of the performance with/without pre-training using Multi-PIE on CK + dataset.

Empty Cell| 2B(N + M) [6]| IDFERM

---|---|---

Empty Cell| Without Pre-training| With Pre-training| Without Pre-training| With Pre-training

CK + (7-class)| 97.03%| 97.10%| 98.32%| 98.35%

MMI| 78.46%| 78.53%| 81.11%| 81.13%

****MMI Dataset**** [40]: This dataset consists of 213 sequences, 208 sequences from this data set containing frontal-view faces of 31 subjects were used in our experiment as in [6]. Since the actual location of the peak frame is not provided, we collect three frames in the middle of each image sequence and associate them with the labels, which results in 624 images in our experiments as in [6], [11]. We divided the MMI dataset into 10 subsets for person-independent ten-fold cross validation. The sequence-level predictions are obtained by choosing the class with the highest average score of the three images. Consequently, 10-fold cross validation was conducted. This dataset could be suitable to measure the recognition performance in realistic situations when compared to other datasets.

With the identity-disentangled FER representation, the proposed methods achieve substantial improvements over the previous best performance in MMI dataset as shown in Table 3 and Fig. 13(c). The HNG can further boost the accuracy by incorporating the prior information of normalized face and the relationships of expressions within an applicable framework. Note that the image sequences in the MMI dataset contain a full temporal pattern of expressions, *_i.e._*, from neutral to apex, and then relaxed, and are especially favored by these methods exploiting temporal information.

****Oulu-CASIA VIS Dataset**** [41]: This dataset consists of 480 image sequences of 80 individuals. This dataset is captured under the visible (VIS) normal illumination conditions and is a subset of Oulu-CASIA NIR-VIS dataset. Each individual poses six basic expressions as in MMI dataset. Only the last three frames are used for individual-independent 10-fold cross validation, and the total number of images is 1440 as in [5].

In Oulu-CASIA dataset, the IDFERM performs well in recognizing fear and happy expressions, while angry is the hardest expression, which is mostly confused with disgust as shown in Fig. 13(d). The performance results are shown in Table 6 and are similar to those on the CK + and MMI datasets.

Table 6. Performance comparison of the rank-1 recognition accuracy on the Oulu-CASIA VIS dataset in terms of the 6 expressions (without neutral expression).

Methods| Oulu-CASIA VIS

---|---

STM-ExpLet(video) [9]| 74.59%

DTAGN(video) [10]| 81.46%

PPDN [48]| 84.59%

FN2EN [5]| 87.1%

IDFERM| 88.25%

6\.. Conclusions

We proposed and investigated a novel recognition via generation scheme termed IDFERM to disentangle the identity factors from other factors that are

responsible for facial expression. The anchor-selection and threshold-tuning problems present in previous approaches have been addressed in our proposed adaptive deep metric learning paradigm. The identity-preserving neutral face image generation is efficient for hard negative mining which requires fewer similarity comparisons. However, our gray-scale image processing can lead to information loss as the image quality is not emphasized in our framework as in conventional image generation methods. Also, the adversarial game at image-level is usually time-consuming. In future work, we intend to apply some commonly used visual quality assessment methods for the generated images on top of our model for better texture. Recent feature-level GAN's backbone can be utilized to extend our framework for faster, more stable convergence training, and more complex data structure (e.g., color images). We also expect that the application of recognition via generation idea can facilitate several other closely related tasks, e.g., face recognition, person re-ID, and pose-invariant classification.

Recommended articles

References

1. [1]

J.F. Cohn, P. Ekman

"Measuring facial action

The New Handbook of Methods in Nonverbal Behavior Research (2005), pp. 9-64

[Crossref](#)[Google Scholar](#)

2. [2]

P. Ekman, L. Erika

What the Face reveals: Basic and Applied Studies of Spontaneous Expression

Using the Facial Action Coding System (FACS)

Oxford University Press, USA (1997)

[Google Scholar](#)

3. [3]

J. Zhang, J. Yu, D. Tao

Local deep-feature alignment for unsupervised dimension reduction

IEEE Trans. Image Process., 27 (5) (2018), pp. 2420-2432

[Crossref](#)[View in Scopus](#)[Google Scholar](#)

4. [4]

S. Han, Z. Meng, A. Khan, Y. Tong

Incremental boosting convolutional neural network for facial action unit recognition

Advances in Neural Information Processing Systems (2016), pp. 109-117

[View in Scopus](#)[Google Scholar](#)

5. [5]

H. Ding, S. Zhou, R. Chellappa

Facenet2expnet: regularizing a deep face recognition net for expression recognition

Automatic Face & Gesture Recognition (FG 2017), 2017 12th IEEE International Conference on, IEEE (2017), pp. 118-126

[View in Scopus](#)[Google Scholar](#)

6. [6]

X. Liu, B.V.K. Vijaya Kumar, J. You, P. Jia

Adaptive deep metric learning for identity-aware facial expression recognition

Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (2017), pp. 522-531

[View in Scopus](#)[Google Scholar](#)

7. [7]

X. Liu, B.V.K. Vijaya Kumar, Y. Ge, C. Yang, J. You, P. Jia

Normalized face generation via perceptual generative adversarial networks
Identity, Security, and Behavior Analysis (ISBA), 2018 IEEE 4th International
Conference on, IEEE (2018), pp. 1-8

[View PDF](#)[View article](#)[Google Scholar](#)

8. [8]

J. Haxby, E. Hoffman, M. Gobbini

"The distributed human neural system for face perception

Trends Cogn. Sci., 4 (6) (2000), pp. 223-233

[View PDF](#)[View article](#)[View in Scopus](#)[Google Scholar](#)

9. [9]

M. Liu, S. Shan, R. Wang, X. Chen

Learning expressionlets on spatio-temporal manifold for dynamic facial
expression recognition

Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition
(2014), pp. 1749-1756

[View in Scopus](#)[Google Scholar](#)

10. [10]

H. Jung, S. Lee, J. Yim, S. Park, J. Kim

Joint fine-tuning in deep neural networks for facial expression recognition

Proceedings of the IEEE International Conference on Computer Vision (2015),
pp. 2983-2991

[View in Scopus](#)[Google Scholar](#)

11. [11]

Z. Meng, P. Liu, J. Cai, Y. Tong

Identity-aware convolutional neural network for facial expression recognition

Automatic Face & Gesture Recognition (FG 2017), 2017 12th IEEE International
Conference on, IEEE (2017), pp. 558-565

[View in Scopus](#)[Google Scholar](#)

12. [12]

E. Sariyanidi, H. Gunes, A. Cavallaro

Automatic analysis of facial affect: a survey of registration, representation
and recognition

IEEE Trans. Pattern Anal. Mach. Intell., 37 (2015), pp. 1113-1133

[View in Scopus](#)[Google Scholar](#)

13. [13]

S. Rifai, Y. Bengio, A. Courville

Disentangling factors of variation for facial expression recognition

Computer Vision?ECCV 2012, Berlin, Heidelberg, Springer (2012), pp. 808-822

[Crossref](#)[View in Scopus](#)[Google Scholar](#)

14. [14]

P. Liu, J. Zhou, I. Tsang, Z. Meng, S. Han, Y. Tong

Feature disentangling machine-a novel approach of feature selection and
disentangling in facial expression analysis

European Conference on Computer Vision, Cham, Springer (2014), pp. 151-166

[Crossref](#)[View in Scopus](#)[Google Scholar](#)

15. [15]

H. Liu, Y. Tian, Y. Yang, L. Pang, T. Huang

Deep relative distance learning: tell the difference between similar vehicles

Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition
(2016), pp. 2167-2175

[View in Scopus](#)[Google Scholar](#)

16. [16]

W. Deng, J. Hu, Z. Wu, J. Guo

From one to many: pose-aware metric learning for single-sample face

recognition

Pattern Recognit., 77 (2018), pp. 426-437

[View PDF](#)[View article](#)[View in Scopus](#)[Google Scholar](#)

17. [17]

J. Wang, Z. Wang, C. Liang, C. Gao, N. Sang

Equidistance constrained metric learning for person re-identification

Pattern Recognit., 74 (2018), pp. 38-51

[View PDF](#)[View article](#)[Crossref](#)[Google Scholar](#)

18. [18]

S. Chopra, R. Hadsell, Y. LeCun

Learning a similarity metric discriminatively, with application to face verification

Computer Vision and Pattern Recognition, 2005. CVPR 2005. IEEE Computer Society Conference on, 1, IEEE (2005), pp. 539-546

[Crossref](#)[View in Scopus](#)[Google Scholar](#)

19. [19]

S. Ding, L. Lin, G. Wang, H. Chao

Deep feature learning with relative distance comparison for Person _Re_ identification

Pattern Recognit., 48 (2015), pp. 2993-3003

[View PDF](#)[View article](#)[View in Scopus](#)[Google Scholar](#)

20. [20]

K. Sohn

Improved deep metric learning with multi-class n-pair loss objective

Advances in Neural Information Processing Systems (2016), pp. 1857-1865

[View in Scopus](#)[Google Scholar](#)

21. [21]

Y. Dong, B. Du, L. Zhang, L. Zhang, D. Tao

LAM3L: Locally adaptive maximum margin metric learning for visual data classification

Neurocomputing, 235 (2017), pp. 1-9

[View PDF](#)[View article](#)[Crossref](#)[Google Scholar](#)

22. [22]

S. Thrun, others

Robotic mapping: A survey

Exploring Artif. Intell. New Millennium, 1 (1-35) (2002), p. 1

[Google Scholar](#)

23. [23]

Christopher M. Bishop

Pattern recognition and machine learning (information science and statistics) springer-verlag new york

Inc. Secaucus, NJ, USA (2006)

[Google Scholar](#)

24. [24]

Sebastian Thrun

Robotic mapping: A survey

Explor. Artif. Intell. New Millenn., 1 (2002), pp. 1-35

[Google Scholar](#)

25. [25]

P. Vincent, H. Larochelle, Y. Bengio

Extracting and composing robust features with denoising autoencoders

Proceedings of the 25th international conference on Machine learning, ACM (2008), pp. 1096-1103

[Crossref](#)[Google Scholar](#)

26. [26]
I. Goodfellow, J. Pouget, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, Y. Bengio
Generative adversarial nets
Advances in neural information processing systems (2014), pp. 2672-2680
Google Scholar
27. [27]
I. Goodfellow, Ian. "NIPS 2016 tutorial: Generative adversarial networks." In arXiv: 1701.00160, 2016.
Google Scholar
28. [28]
L. Tran, X. Yin, X. Liu
Disentangled representation learning gan for pose-invariant face recognition
Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 3 (2017), p. 7
Google Scholar
29. [29]
P. Omkar, A. Vedaldi, A. Zisserman
Deep Face Recognition
BMVC, 1 (3) (2015), p. 6
Google Scholar
30. [30]
X. Wu, R. He, Z. Sun, T. Tan
A Light CNN for Deep Face Representation with Noisy Labels
IEEE Trans. Inf. Forensics Secur., 13 (11) (2018), pp. 2884-2896
Google Scholar
31. [31]
L. Theis, A. van den Oord, and M. Bethge. "A note on the evaluation of generative models." arXiv preprint arXiv:1511.01844 (2015).
Google Scholar
32. [32]
Y. Wen, K. Zhang, Z. Li, Y. Qiao
A discriminative feature learning approach for deep face recognition
European Conference on Computer Vision, Cham, Springer (2016), pp. 499-515
CrossrefView in ScopusGoogle Scholar
33. [33]
E. Barsoum, C. Zhang, C.C. Ferrer, Z. Zhang
Training deep networks for facial expression recognition with crowd-sourced label distribution
Proceedings of the 18th ACM International Conference on Multimodal Interaction, ACM (2016), pp. 279-283
CrossrefView in ScopusGoogle Scholar
34. [34]
H. Shi, Y. Yang, X. Zhu, L. Zhen, W. Zheng, S.Z. Li.
Embedding deep metric for person re-identification: a study against large variations
European Conference on Computer Vision, Cham, Springer (2016), pp. 732-748
CrossrefView in ScopusGoogle Scholar
35. [35]
Y. Sun, Y. Chen, X. Wang, X. Tang
Deep learning face representation by joint identification-verification
Advances in Neural Information Processing Systems (2014), pp. 1988-1996
View in ScopusGoogle Scholar
36. [36]

R. Gross, I. Matthews, J. Cohn, T. Kanade, S. Baker

Multi-pie

Image Vision Comput., 28 (5) (2010), pp. 807-813

[View PDF](#)[View article](#)[View in Scopus](#)[Google Scholar](#)

37. [37]

K. He, X. Zhang, S. Ren, J. Sun

Delving deep into rectifiers: surpassing human-level performance on imagenet classification

Proceedings of the IEEE International Conference on Computer Vision (2015),

pp. 1026-1034

[Crossref](#)[View in Scopus](#)[Google Scholar](#)

38. [38]

M. Abadi, P. Barham, J. Chen, Z. Chen, A. Davis, J. Dean, M. Devin, S.

Ghemawat, G. Irving, M. Isard, M. Kudlur, J. Levenberg, R. Monga, S. Moore,

D.G. Murray, B. Steiner, P. Tucker, V. Vasudevan, P. Warden, M. Wicke, Y. Yu,

X. Zheng

TensorFlow: a system for large-scale machine learning

OSDI, 16 (2016), pp. 265-283

[Google Scholar](#)

39. [39]

P. Lucey, J.F. Cohn, T. Kanade, J. Saragih, Z. Ambadar, I. Matthews

The extended cohn-kanade dataset (ck+): A complete dataset for action unit and emotion-specified expression

Computer Vision and Pattern Recognition Workshops (CVPRW), 2010 IEEE Computer Society Conference on, IEEE (2010), pp. 94-101

[View in Scopus](#)[Google Scholar](#)

40. [40]

M. Pantic, M. Valstar, R. Rademaker, L. Maat

Web-based database for facial expression analysis

2005 IEEE international conference on multimedia and Expo, IEEE (2005), p. 5

[Google Scholar](#)

41. [41]

G. Zhao, X. Huang, M. Taini, S.Z. Li, M. Pietikainen

Facial expression recognition from near-infrared videos

Image Vision Comput., 29 (9) (2011), pp. 607-619

[View PDF](#)[View article](#)[View in Scopus](#)[Google Scholar](#)

42. [42]

P. Liu, S. Han, Z. Meng, Y. Tong

Facial expression recognition via a boosted deep belief network

Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (2014), pp. 1805-1812

[View in Scopus](#)[Google Scholar](#)

43. [43]

W. Chu, F. Torre, J.F. Cohn

Selective transfer machine for personalized facial expression analysis

IEEE Trans. Pattern Anal. Mach. Intell., 39 (3) (2017), pp. 529-545

[View in Scopus](#)[Google Scholar](#)

44. [44]

S. Lee, W. Baddar, Y. Ro.

Collaborative expression representation using peak expression and intra class variation face images for practical subject-independent emotion recognition in videos

Pattern Recognit., 54 (2016), pp. 52-67

[View PDF](#)[View article](#)[View in Scopus](#)[Google Scholar](#)

45. [45]

H. Yan

Collaborative discriminative multi-metric learning for facial expression recognition in video

Pattern Recognit., 75 (2018), pp. 33-40

[View PDF](#)[View article](#)[Google Scholar](#)

46. [46]

M. Liu, S. Li, S. Shan, X. Chen

Au-inspired deep net-works for facial expression feature learning

Neurocomputing, 159 (2015), pp. 126-136

[View PDF](#)[View article](#)[View in Scopus](#)[Google Scholar](#)

47. [47]

P. Khorrami, T. Paine, T. Huang

Do deep neural networks learn facial action units when doing expression recognition?

Proceedings of the IEEE International Conference on Computer Vision Workshops (2015), pp. 19-27

[View in Scopus](#)[Google Scholar](#)

48. [48]

X. Zhao, X. Liang, L. Liu, T. Li, N. Vasconcelos, S. Yan

Peakpiloted deep network for facial expression recognition

European conference on computer vision, Cham, Springer (2016), pp. 425-442

[Crossref](#)[View in Scopus](#)[Google Scholar](#)

Cited by (86)

* ### A systematic review on affective computing: emotion models, databases, and recent advances 2022, Information Fusion

[Show abstract](#)

Affective computing conjoins the research topics of emotion recognition and sentiment analysis, and can be realized with unimodal or multimodal data, consisting primarily of physical information (e.g., text, audio, and visual) and physiological signals (e.g., EEG and ECG). Physical-based affect recognition caters to more researchers due to the availability of multiple public databases, but it is challenging to reveal one's inner emotion hidden purposefully from facial expressions, audio tones, body gestures, etc. Physiological signals can generate more precise and reliable emotional results; yet, the difficulty in acquiring these signals hinders their practical application. Besides, by fusing physical information and physiological signals, useful features of emotional states can be obtained to enhance the performance of affective computing models. While existing reviews focus on one specific aspect of affective computing, we provide a systematical survey of important components: emotion models, databases, and recent advances. Firstly, we introduce two typical emotion models followed by five kinds of commonly used databases for affective computing. Next, we survey and taxonomize state-of-the-art unimodal affect recognition and multimodal affective analysis in terms of their detailed architectures and performances. Finally, we discuss some critical aspects of affective computing and its applications and conclude this review by pointing out some of the most promising future directions, such as the establishment of benchmark database and fusion strategies. The overarching goal of this systematic review is to help academic and industrial researchers understand the recent advances as well as new developments in this fast-paced, high-impact domain.

* ### OAENet: Oriented attention ensemble for accurate facial expression recognition 2021, Pattern Recognition

[Show abstract](#)

Facial Expression Recognition (FER) is a challenging yet important research topic owing to its significance with respect to its academic and commercial potentials. In this work, we propose an oriented attention pseudo-siamese network that takes advantage of global and local facial information for high accurate FER. Our network consists of two branches, a maintenance branch that consisted of several convolutional blocks to take advantage of high-level semantic features, and an attention branch that possesses a UNet-like architecture to obtain local highlight information. Specifically, we first input the face image into the maintenance branch. For the attention branch, we calculate the correlation coefficient between a face and its sub-regions. Next, we construct a weighted mask by correlating the facial landmarks and the correlation coefficients. Then, the weighted mask is sent to the attention branch. Finally, the two branches are fused to output the classification results. As such, a direction-dependent attention mechanism is established to remedy the limitation of insufficient utilization of local information. With the help of our attention mechanism, our network not only grabs a global picture but can also concentrate on important local areas. Experiments are carried out on 4 leading facial expression datasets. Our method has achieved a very appealing performance compared to other state-of-the-art methods.

* ### Sparse deep feature learning for facial expression recognition

2019, Pattern Recognition

Show abstract

While weight sparseness-based regularization has been used to learn better deep features for image recognition problems, it introduced a large number of variables for optimization and can easily converge to a local optimum. The L2-norm regularization proposed for face recognition reduces the impact of the noisy information, while expression information is also suppressed during the regularization. A feature sparseness-based regularization that learns deep features with better generalization capability is proposed in this paper. The regularization is integrated into the loss function and optimized with a deep metric learning framework. Through a toy example, it is showed that a simple network with the proposed sparseness outperforms the one with the L2-norm regularization. Furthermore, the proposed approach achieved competitive performances on four publicly available datasets, i.e., FER2013, CK+, Oulu-CASIA and MMI. The state-of-the-art cross-database performances also justify the generalization capability of the proposed approach.

* ### Deep Facial Expression Recognition: A Survey

2022, IEEE Transactions on Affective Computing

* ### Facial Expression Recognition of Industrial Internet of Things by Parallel Neural Networks Combining Texture Features

2021, IEEE Transactions on Industrial Informatics

* ### Fine-Grained Facial Expression Recognition in the Wild

2021, IEEE Transactions on Information Forensics and Security

View all citing articles on Scopus

****Xiaofeng Liu**** received the B.Eng. degree in automation and B.A. degree in communication from the University of Science and Technology of China, Hefei, China, in 2014. He was the research assistant in MSRA, CMU and NTU. He is currently pursuing his Ph.D. at the University of Chinese Academy of Sciences, Beijing and a Research Associate in the Department of Electrical and Computer Engineering, Carnegie Mellon University, Pittsburgh. He was a recipient of the Best Paper award of the IEEE International Conference on Identity, Security and Behavior Analysis 2018. His research interests include image processing, computer vision, and pattern recognition.

****B.V.K. Vijaya Kumar**** is the U.A. & Helen Whitaker Professor of Electrical

and Computer Engineering at Carnegie Mellon University, Pittsburgh. He is also the Director of Carnegie Mellon University Africa in Rwanda. His research interests include computer vision, pattern recognition algorithms, and applications, coding, and signal processing for data storage systems. His publications include a book entitled Correlation Pattern Recognition, 22 book chapters, more than 400 conference papers, and more than 200 journal papers. He is also a co-inventor of 12 patents. He served as a topical editor of *_Applied Optics_* and as an associate editor of the *IEEE _Transactions on Information Forensics and Security_*. He has served on many conference program committees and was co-chair of several conference program committees. In 2003, he received the Eta Kappa Nu Award for Excellence in Teaching in the ECE Department, CMU and the Carnegie Institute of Technology's Dowd Fellowship for educational contributions and he was a co-recipient of the 2008 Outstanding Faculty Research Award in CMU's College of Engineering. He is a fellow of IEEE, SPIE, OSA, AAAS, IAPR and NAI.

****Ping Jia**** received his B.Eng and MSc degree in computer science from the University of Science and Technology of China, Hefei and his Ph.D. from the Graduate University of the Chinese Academy of Sciences. He is the President of Changchun Institute of Optics, Fine Mechanics and Physics, Chinese Academy of Sciences. His current research interests include image processing, computer vision, and optical engineering.

****Jane You**** received the Ph.D. degree from La Trobe University, Melbourne, VIC, Australia, in 1992. She is currently a Professor with the Department of Computing, Hong Kong Polytechnic University, Hong Kong, and the Chair of Department Research Committee. She has researched extensively in the fields of image processing, medical imaging, computer-aided diagnosis, and pattern recognition. She has been a Principal Investigator for one ITF project, three GRF projects, and many other joint grants since she joined PolyU in 1998. Prof. You was a recipient of three awards including Hong Kong Government Industrial Awards, the Special Prize and Gold Medal with Jury's Commendation at the 39th International Exhibition of Inventions of Geneva in 2011 for her current work on retinal imaging, and the Second Place in an International Competition [SPIE Medical Imaging?2009 Retinopathy Online Challenge in (ROC?2009)]. Her research output on retinal imaging has been successfully led to technology transfer with clinical applications. She is an Associate Editor of Pattern Recognition and other journals.

[View Abstract](#)

© 2018 Elsevier Ltd. All rights reserved.

Recommended articles

*** ### An effective deep network using target vector update modules for image restoration**
Pattern Recognition, Volume 122, 2022, Article 108333

Sen Zhai, ?, Linbo Qing

[View PDF](#)

*** ### Robust Low-rank subspace segmentation with finite mixture noise**
Pattern Recognition, Volume 93, 2019, pp. 55-67

Xianglin Guo, ?, Jun Wang

[View PDF](#)

*** ### Noise-robust dictionary learning with slack block-Diagonal structure for face recognition**
Pattern Recognition, Volume 100, 2020, Article 107118

Zhe Chen, ?, Josef Kittler

[View PDF](#)

*** ### Learning disentangling and fusing networks for face completion under structured occlusions**
Pattern Recognition, Volume 99, 2020, Article 107073

Zhihang Li, ?, Zhenan Sun

[View PDF](#)

* ### Strengthening mechanism of lightweight cellular concrete filled with fly ash
Construction and Building Materials, Volume 251, 2020, Article 118954

Xin Liu, ?, Lizhi Sun

[View PDF](#)

* ### Surrogate network-based sparseness hyper-parameter optimization for deep expression recognition
Pattern Recognition, Volume 111, 2021, Article 107701

Weicheng Xie, ?, Meng Yang

[View PDF](#)

[Show 3 more articles](#)

[## Article Metrics](#)

[Citations](#)

* Citation Indexes: 86

[Captures](#)

* Readers: 62

[View details](#)

* [About ScienceDirect](#)

* [Remote access](#)

* [Shopping cart](#)

* [Advertise](#)

* [Contact and support](#)

* [Terms and conditions](#)

* [Privacy policy](#)

[Cookies are used by this site. Cookie Settings](#)

All content on this site: Copyright © 2024 Elsevier B.V., its licensors, and contributors. All rights are reserved, including those for text and data mining, AI training, and similar technologies. For all open access content, the Creative Commons licensing terms apply.

[## Cookie Preference Center](#)

We use cookies which are necessary to make our site work. We may also use additional cookies to analyse, improve and personalise our content and your digital experience. For more information, see our [Cookie Policy](#) and the list of [Google Ad-Tech Vendors](#).

You may choose not to allow some types of cookies. However, blocking some types may impact your experience of our site and the services we are able to offer. See the different category headings below to find out more or change your settings.

[Allow all](#)

[### Manage Consent Preferences](#)

[#### Strictly Necessary Cookies](#)

[Always active](#)

These cookies are necessary for the website to function and cannot be switched off in our systems. They are usually only set in response to actions made by you which amount to a request for services, such as setting your privacy preferences, logging in or filling in forms. You can set your browser to block or alert you about these cookies, but some parts of the site will not then work. These cookies do not store any personally identifiable information.

[Cookie Details List?](#)

[#### Functional Cookies](#)

[Functional Cookies](#)

These cookies enable the website to provide enhanced functionality and personalisation. They may be set by us or by third party providers whose services we have added to our pages. If you do not allow these cookies then some or all of these services may not function properly.

Cookie Details List?

Performance Cookies

Performance Cookies

These cookies allow us to count visits and traffic sources so we can measure and improve the performance of our site. They help us to know which pages are the most and least popular and see how visitors move around the site.

Cookie Details List?

Targeting Cookies

Targeting Cookies

These cookies may be set through our site by our advertising partners. They may be used by those companies to build a profile of your interests and show you relevant adverts on other sites. If you do not allow these cookies, you will experience less targeted advertising.

Cookie Details List?

Back Button

Cookie List

Search Icon

Filter Icon

Clear

checkbox label label

Apply Cancel

Consent Leg.Interest

checkbox label label

checkbox label label

checkbox label label

Confirm my choices